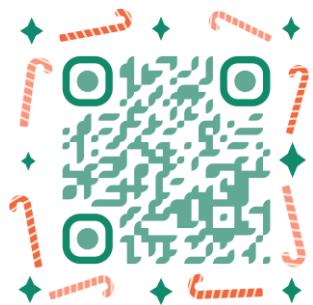


From Multimodal LLM to Human-level AI

Modality, Instruction, Reasoning, Efficiency and Beyond

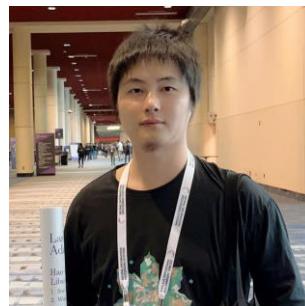


<https://mllm2024.github.io/CVPR2024/>



CC BY 4.0 DEED

Attribution 4.0 International



Hao Fei

National University of Singapore



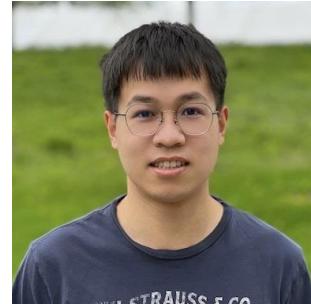
Yuan Yao

National University of Singapore



Ao Zhang

National University of Singapore



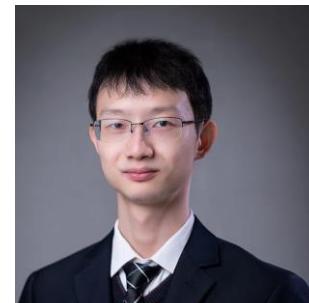
Haotian Liu

University of Wisconsin-Madison



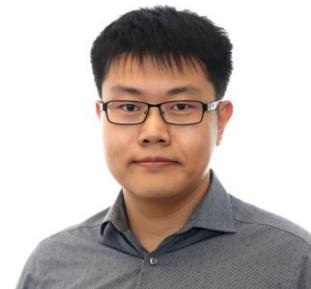
Fuxiao Liu

University of Maryland, College Park



Zhuosheng Zhang

Shanghai Jiao Tong University



Hanwang Zhang

Nanyang Technological University



Shuicheng Yan

Kunlun 2050 Research, Skywork AI

✳ Part-III

Modality and Functionality

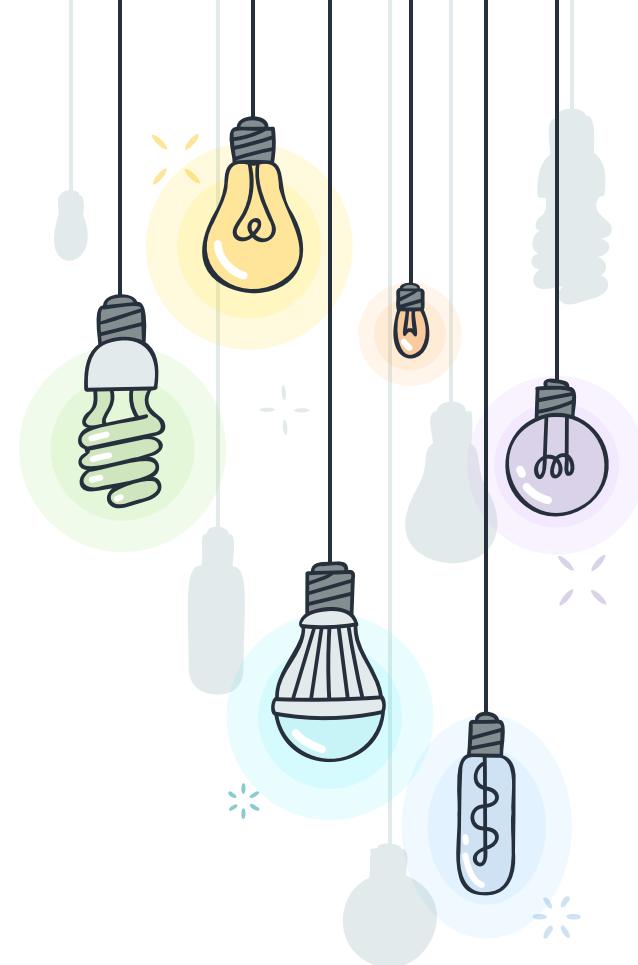


Hao Fei

Research Fellow

National University of Singapore

<http://haofei.vip/>



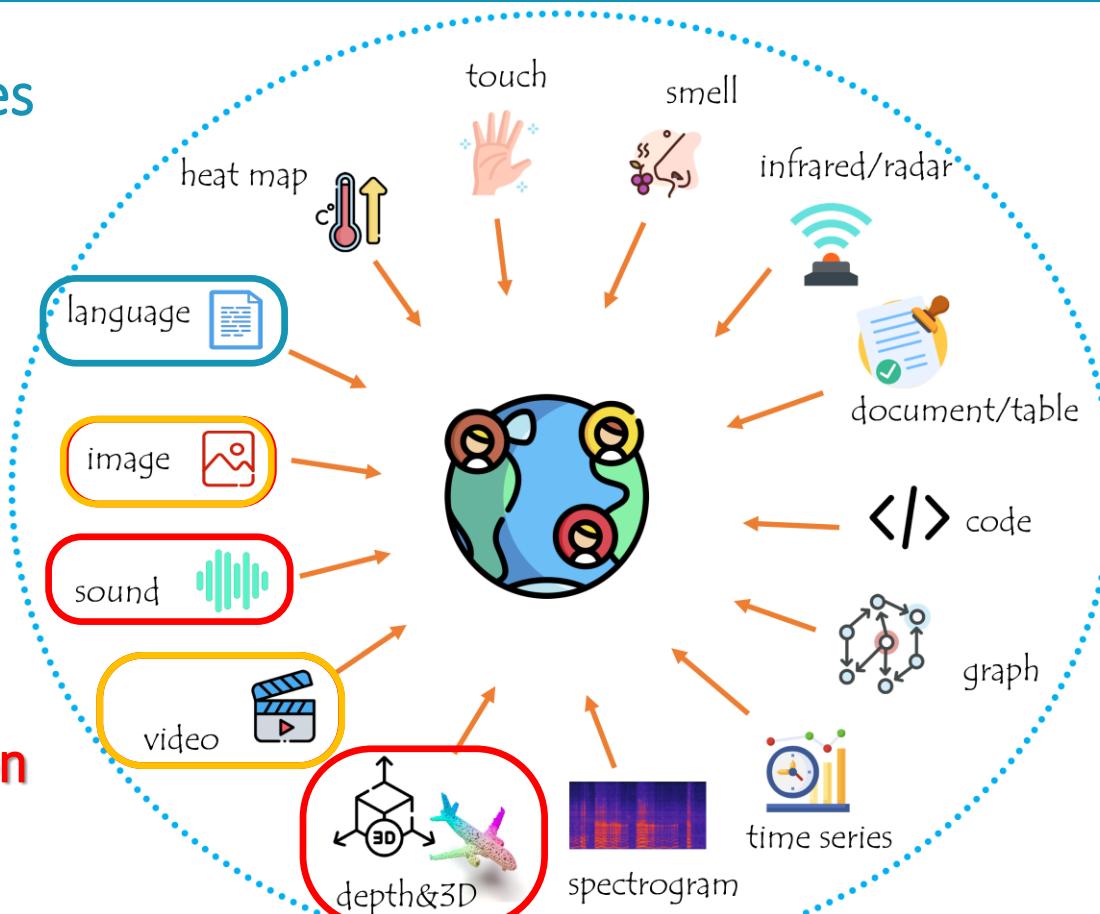
✳️ Table of Content

+ Modality & Functionality

- ✖ Overview
- ✖ Multimodal Perceiving
- ✖ Multimodal Generation
- ✖ Unified MLLM
- ✖ Fine-grained MLLM
- ✖ What's Next

Overview of Modality and Functionality

- **Modalities**



Language + Vision

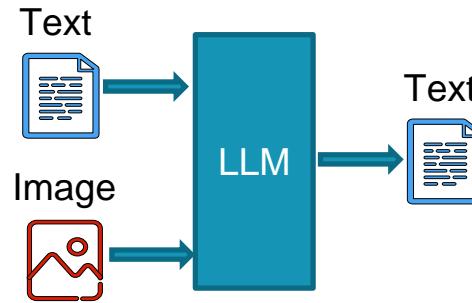
✳️ Overview of Modality and Functionality

	Modality (w/ Language)			
	Image	Video	Audio	3D
Input-side Perceiving	Flamingo, Kosmos-1, Blip2, mPLUG-Owl, Mini-GPT4, LLaVA, InstructBLIP, VPGTrans, CogVLM, Monkey, Chameleon, Otter, Qwen-VL, GPT-4v, SPHINX, Yi-VL, Fuyu, ...	VideoChat, Video-ChatGPT, Video-LLaMA, PandaGPT, MovieChat, Video-LLaVA, LLaMA-VID, Momentor, ...	AudioGPT, SpeechGPT, VIOLA, AudioPaLM, SALMONN, MU-LLaMA, ...	3D-LLM, 3D-GPT, LL3DA, SpatialVLM, PointLLM, PointBind, ...
	[Pixel-wise] GPT4RoI, LION, MiniGPT-v2, NExT-Chat, Kosmos-2, GLaMM, LISA, DetGPT, Osprey, PixelLM, ...	[Pixel-wise] PG-Video-LLaVA, Merlin, MotionEpic, ...	-	-
	Video-LLaVA, Chat-UniVi, LLaMA-VID		-	-
	Panda-GPT, Video-LLaMA, AnyMAL, Macaw-LLM, Gemini, VideoPoet, ImageBind-LLM, LLMBind, LLaMA-Adapter, ...			-
Perceiving + Generating	GILL, EMU, MiniGPT-5, DreamLLM, LLaVA-Plus, InternLM-XComposer2, SEED-LLaMA, LaVIT, Mini-Gemini, ...	GPT4Video, Video-LaVIT, VideoPoet, ...	AudioGPT, SpeechGPT, VIOLA, AudioPaLM, ...	-
	[Pixel-wise] Vitron		-	-
	NExT-GPT, Unified-IO 2, AnyGPT, CoDi-2, Modaverse, ViT-Lens, ...			-

Multimodal Perceiving

- **Image-perceiving MLLM**

- + Flamingo,
- + Kosmos-1,
- + Blip2, mPLUG-Owl,
- + Mini-GPT4, LLaVA,
- + InstructBLIP, Otter,
- + VPGTrans
- + Chameleon,
- + Qwen-VL, GPT-4v,
- + SPHINX,
- + ...



Encode input images with external image encoders, generating LLM-understandable visual feature, which is then fed into the LLM. LLM then interprets the input images based on the input text instructions and produces a textual response.

[1] Flamingo: a Visual Language Model for Few-Shot Learning. 2022

[2] Language Is Not All You Need: Aligning Perception with Language Models. 2023

[3] BLIP-2: Bootstrapping Language-Image Pre-training with Frozen Image Encoders and Large Language Models. 2023

[4] MiniGPT-4: Enhancing Vision-Language Understanding with Advanced Large Language Models. 2024

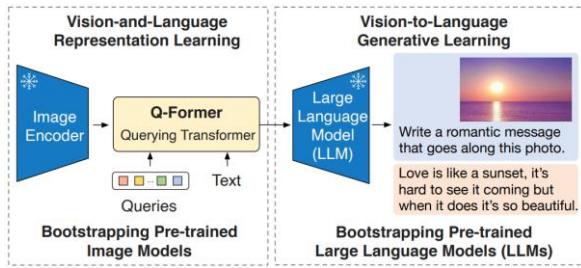
...

Multimodal Perceiving

- Image-perceiving MLLM

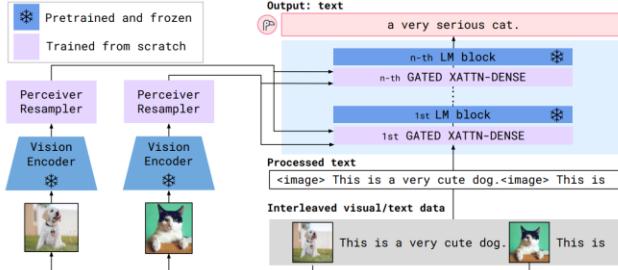
+

Blip2



+

Flamingo



[1] Flamingo: a Visual Language Model for Few-Shot Learning. 2022

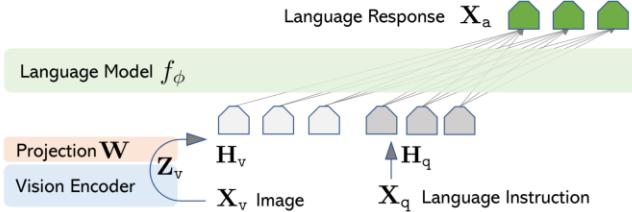
[2] BLIP-2: Bootstrapping Language-Image Pre-training with Frozen Image Encoders and Large Language Models. 2023

[3] Visual Instruction Tuning. 2023

[4] A Survey on Multimodal Large Language Models. <https://github.com/BradyFU/Awesome-Multimodal-Large-Language-Models>, 2023.

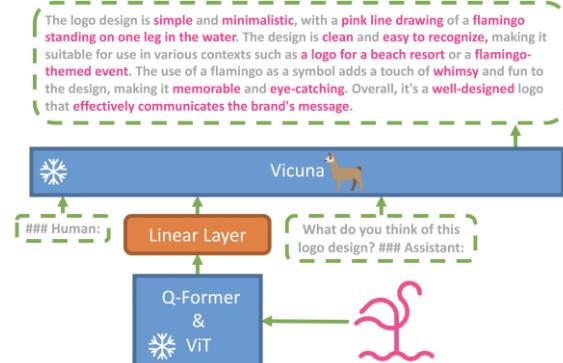
+

LLaVA



+

Mini-GPT4



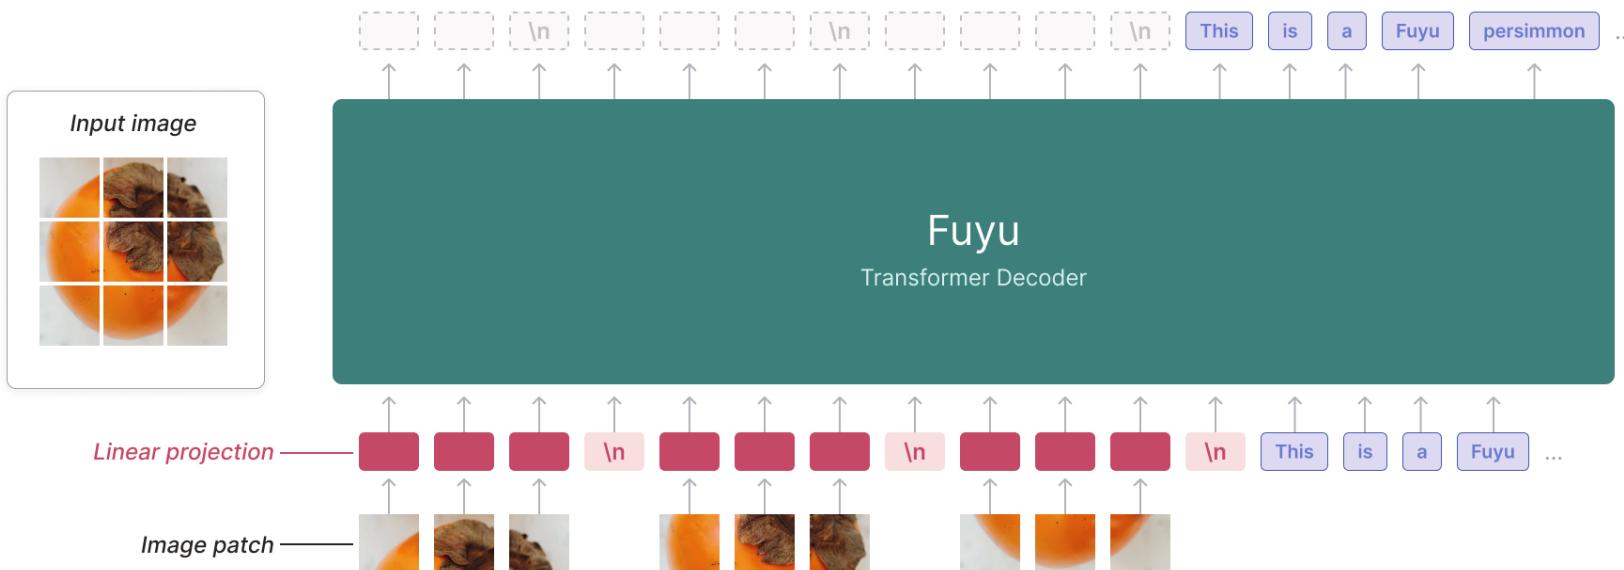
Multimodal Perceiving

- **Image-perceiving MLLM**

- + **Fuyu**



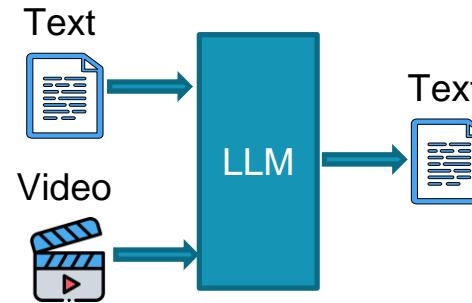
*Unlike all other existing image-oriented MLLMs, Fuyu processes image information without a frontend image encoder, and instead **directly inputs image patches into the LLM for interpretation**.*



Multimodal Perceiving

- Video-perceiving MLLM

- + VideoChat,
- + Video-ChatGPT,
- + Video-LLaMA,
- + PandaGPT,
- + MovieChat,
- + Video-LLaVA,
- + LLaMA-VID,
- + Momentor
- + ...



Encode input videos with external video encoders, generating LLM-understandable visual feature, feeding into LLM, which then interprets the input videos based on the input text instructions and produces a textual response.

[1] VideoChat: Chat-Centric Video Understanding. 2023

[2] Video-ChatGPT: Towards Detailed Video Understanding via Large Vision and Language Models. 2023

[3] Video-LLaMA: An Instruction-tuned Audio-Visual Language Model for Video Understanding. 2023

[4] Video-LLaVA: Learning United Visual Representation by Alignment Before Projection. 2023

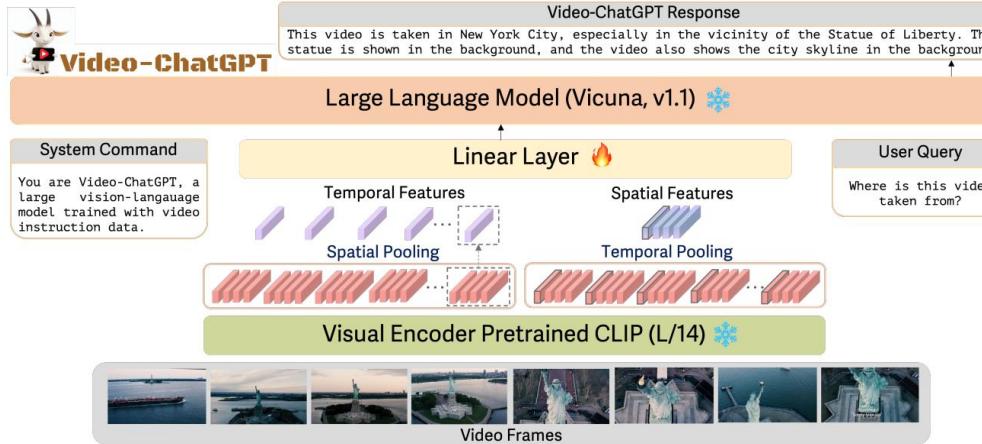
[5] Momentor: Advancing Video Large Language Model with Fine-Grained Temporal Reasoning. 2024

...

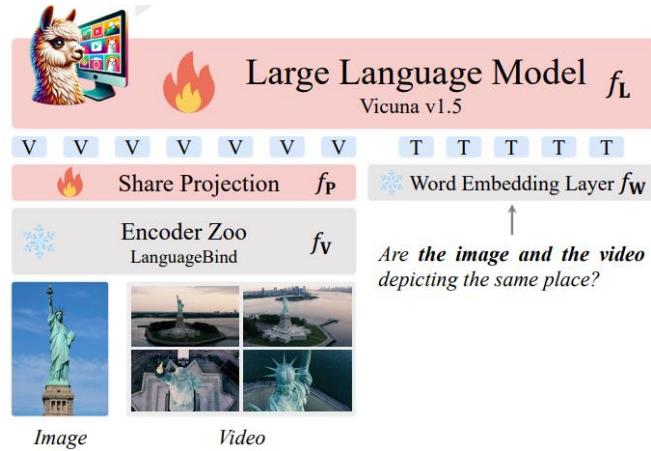
Multimodal Perceiving

- Video-perceiving MLLM

- Video-ChatGPT



- Video-LLaVA



[1] Video-ChatGPT: Towards Detailed Video Understanding via Large Vision and Language Models. 2023

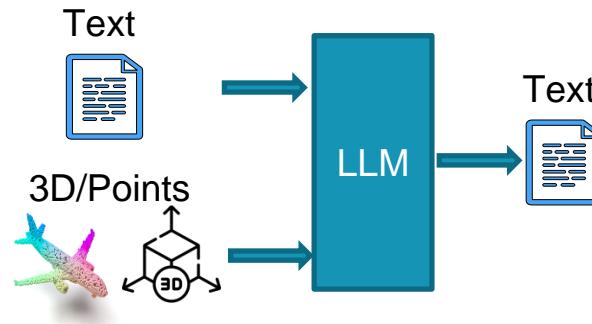
[2] Video-LLaVA: Learning United Visual Representation by Alignment Before Projection. 2023

[3] Video Understanding with Large Language Models: A Survey. <https://github.com/yunlong10/Awesome-LLMs-for-Video-Understanding>, 2023

Multimodal Perceiving

- 3D-perceiving MLLM

- + 3D-LLM,
- + 3D-GPT,
- + LL3DA,
- + SpatialVLM
- + PointLLM
- + Point-Bind
- + ...



Encode input 3D information with external encoders, generating LLM-understandable 3D feature, feeding into LLM, which then interprets the input 3D/points based on the input text instructions and produces a textual response.

[1] 3D-LLM: Injecting the 3D World into Large Language Models. 2023

[2] 3D-GPT: Procedural 3D Modeling with Large Language Models. 2023

[3] LL3DA: Visual Interactive Instruction Tuning for Omni-3D Understanding, Reasoning, and Planning. 2023

[4] PointLLM: Empowering Large Language Models to Understand Point Clouds. 2023

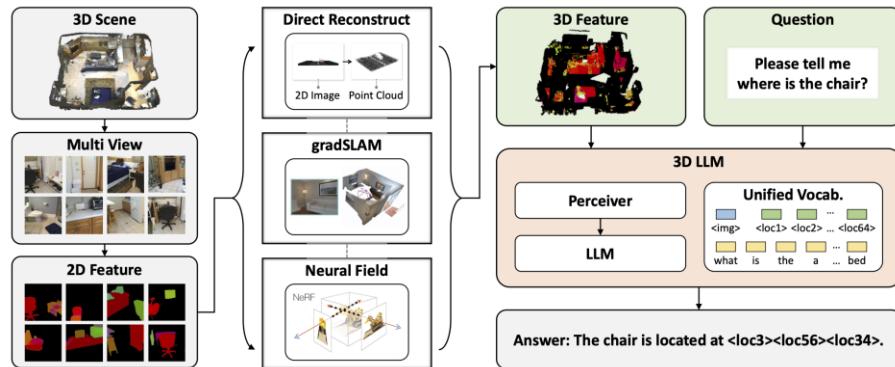
[5] SpatialVLM: Endowing Vision-Language Models with Spatial Reasoning Capabilities. 2024

...

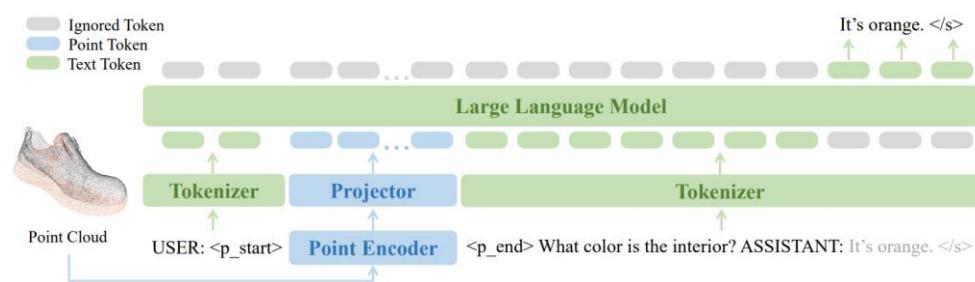
Multimodal Perceiving

- 3D-perceiving MLLM

- + 3D-LLM



- + PointLLM



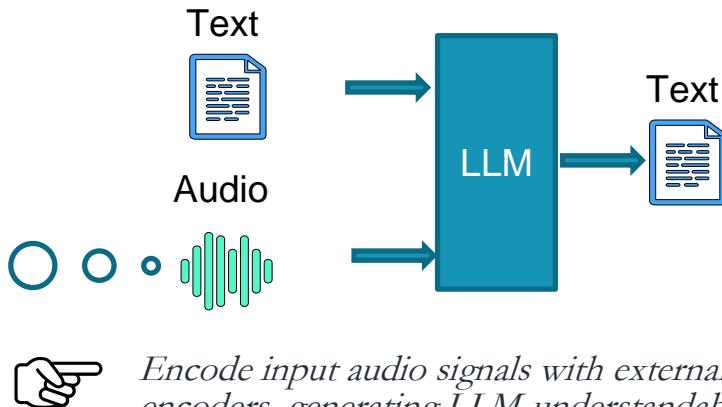
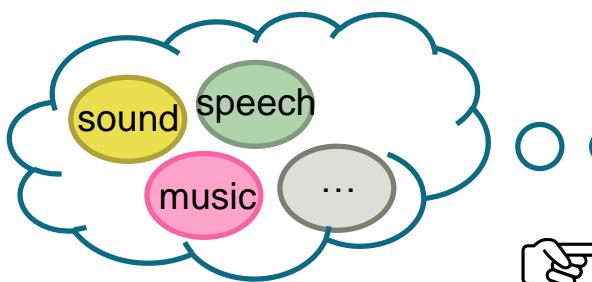
[1] 3D-LLM: Injecting the 3D World into Large Language Models. 2023

[2] PointLLM: Empowering Large Language Models to Understand Point Clouds. 2023

Multimodal Perceiving

- **Audio-perceiving MLLM**

- + AudioGPT,
- + SpeechGPT,
- + VIOLA,
- + AudioPaLM
- + SALMONN
- + MU-LLaMA
- + ...



Encode input audio signals with external encoders, generating LLM-understandable signal features, feeding into LLM, which then interprets the audio based on the input text instructions and produces a textual response.

[1] *AudioGPT: Understanding and Generating Speech, Music, Sound, and Talking Head.* 2023

[2] *SpeechGPT: Empowering Large Language Models with Intrinsic Cross-Modal Conversational Abilities.* 2023

[3] *VioLA: Unified Codec Language Models for Speech Recognition, Synthesis, and Translation.* 2023

[4] *AudioPaLM: A Large Language Model That Can Speak and Listen.* 2023

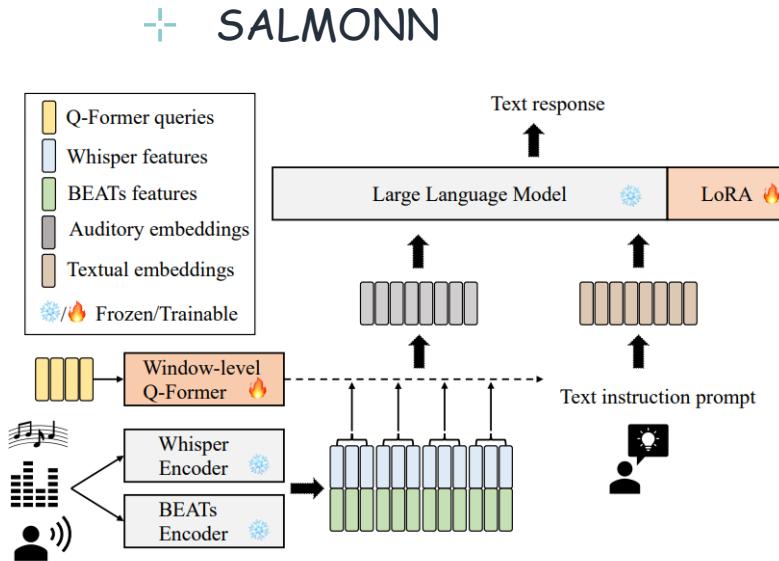
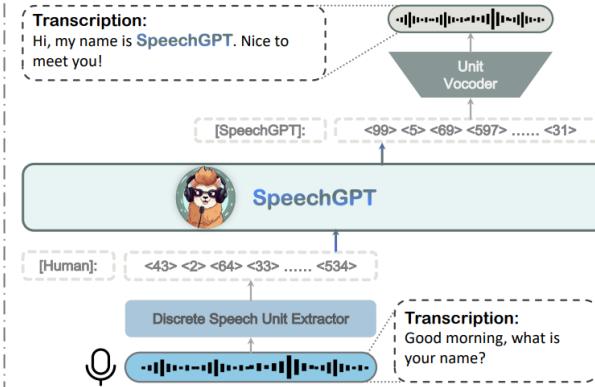
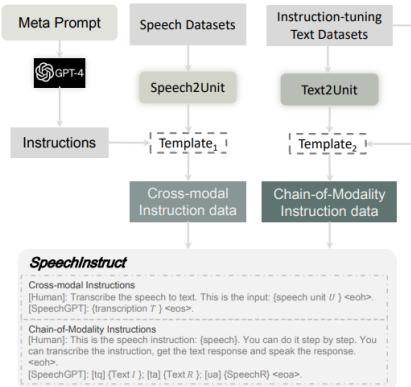
[5] *SALMONN: Towards Generic Hearing Abilities for Large Language Models.* 2023

...

Multimodal Perceiving

- Audio-perceiving MLLM

-  SpeechGPT



[1] *SpeechGPT: Empowering Large Language Models with Intrinsic Cross-Modal Conversational Abilities*. 2023

[2] *SALMONN: Towards Generic Hearing Abilities for Large Language Models*. 2023

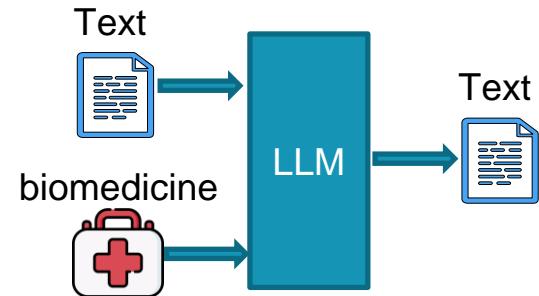
[3] *Sparks of Large Audio Models: A Survey and Outlook*. <https://github.com/EmulationAI/awesome-large-audio-models>, 2023

Multimodal Perceiving

- X-perceiving MLLM

- + Bio-/Medical & Healthcare

+ BioGPT	+ DoctorGLM	+ MedAlpaca
+ DrugGPT	+ BianQue	+ AlpaCare
+ BioMedLM	+ ClinicalGPT	+ Zhongjing
+ OphGLM	+ Qilin-Med	+ PMC-LLaMA
+ GatorTron	+ ChatDoctor	+ CPLLM
+ GatorTronGPT	+ BenTsao	+ MedPaLM 2
+ MEDITRON	+ HuatuoGPT	+ BioMedGPT



[1] BioGPT: Generative Pre-trained Transformer for Biomedical Text Generation and Mining. 2022

[2] DrugGPT: A GPT-based Strategy for Designing Potential Ligands Targeting Specific Proteins. 2023

[3] MEDITRON-70B: Scaling Medical Pretraining for Large Language Models. 2023

[4] HuaTuo: Tuning LLaMA Model with Chinese Medical Knowledge. 2023

[5] AlpaCare: Instruction-tuned Large Language Models for Medical Application. 2023

[6] A Survey of Large Language Models in Medicine: Progress, Application, and Challenge, <https://github.com/AI-in-Health/MedLLMsPracticalGuide>. 2023. 16

Multimodal Perceiving

- X-perceiving MLLM

- + Molecule & Chemistry

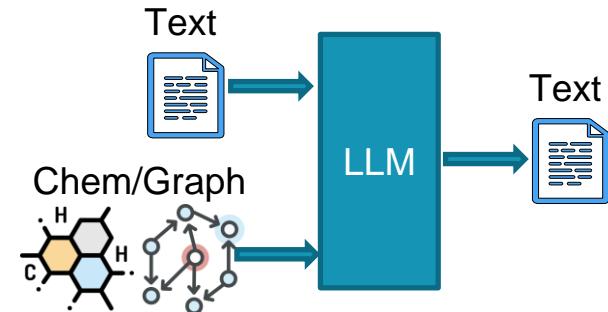
- + ChemGPT
- + SPT
- + T5 Chem
- + ChemLLM
- + MolCA
- + MolXPT
- + MolSTM
- + GIMLET
- + ...

- + Graph

- + StructGPT
- + GPT4Graph
- + GraphGPT
- + LLaGA
- + HiGPT
- + ...

- + Geographical Information System (GIS)

- + GeoGPT



[1] Neural Scaling of Deep Chemical Models. 2022

[2] ChemLLM: A Chemical Large Language Model. 2023

[3] MolCA: Molecular Graph-Language Modeling with Cross-Modal Projector and Uni-Modal Adapter. 2023

[4] StructGPT: A General Framework for Large Language Model to Reason on Structured Data. 2023

[5] LLaGA: Large Language and Graph Assistant. 2023

[6] Awesome-Graph-LLM, <https://github.com/XiaoxinHe/Awesome-Graph-LLM>. 2023

* Unified MLLM: Perceiving + Generation

- Scenarios



*Often, MLLMs need to not only **understand** the input multimodal information, but also to **generate** information in that modality.*

- + Image Captioning
- + Visual Question Answering
- + Text-to-Vision Synthesis
- + Vision-to-Vision Translation
- + Scene Text Recognition
- + Scene Text Inpainting
- + ...

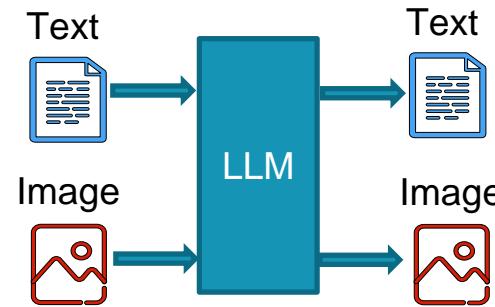
✳️ Overview of Modality and Functionality

	Modality (w/ Language)			
	Image	Video	Audio	3D
Input-side Perceiving	Flamingo, Kosmos-1, Blip2, mPLUG-Owl, Mini-GPT4, LLaVA, InstructBLIP, VPGTrans, CogVLM, Monkey, Chameleon, Otter, Qwen-VL, GPT-4v, SPHINX, Yi-VL, Fuyu, ...	VideoChat, Video-ChatGPT, Video-LLaMA, PandaGPT, MovieChat, Video-LLaVA, LLaMA-VID, Momentor, ...	AudioGPT, SpeechGPT, VIOLA, AudioPaLM, SALMONN, MU-LLaMA, ...	3D-LLM, 3D-GPT, LL3DA, SpatialVLM, PointLLM, PointBind, ...
	[Pixel-wise] GPT4RoI, LION, MiniGPT-v2, NExT-Chat, Kosmos-2, GLaMM, LISA, DetGPT, Osprey, PixelLM, ...	[Pixel-wise] PG-Video-LLaVA, Merlin, MotionEpic, ...	-	-
	Video-LLaVA, Chat-UniVi, LLaMA-VID		-	-
	Panda-GPT, Video-LLaMA, AnyMAL, Macaw-LLM, Gemini, VideoPoet, ImageBind-LLM, LLMBind, LLaMA-Adapter, ...			-
Perceiving + Generating	GILL, EMU, MiniGPT-5, DreamLLM, LLaVA-Plus, InternLM-XComposer2, SEED-LLaMA, LaVIT, Mini-Gemini, ...	GPT4Video, Video-LaVIT, VideoPoet, ...	AudioGPT, SpeechGPT, VIOLA, AudioPaLM, ...	-
	[Pixel-wise] Vitron		-	-
	NExT-GPT, Unified-IO 2, AnyGPT, CoDi-2, Modaverse, ViT-Lens, ...			-

* Unified MLLM: Perceiving + Generation

- **Image**

- + GILL
- + EMU
- + MiniGPT-5
- + DreamLLM
- + LLaVA-Plus
- + LaVIT
- + ...



Central LLMs take as input both texts and images, after semantics comprehension, and generate both texts and images.

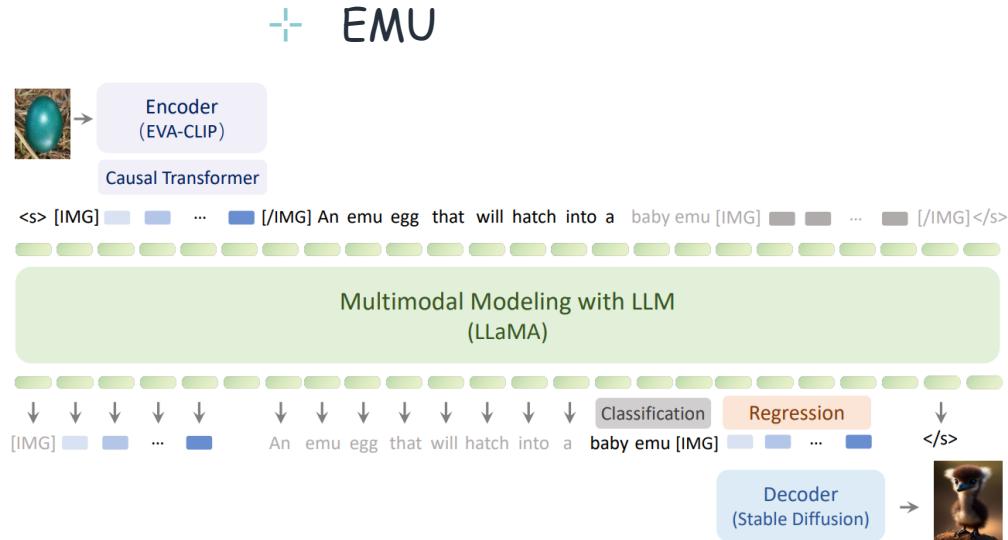
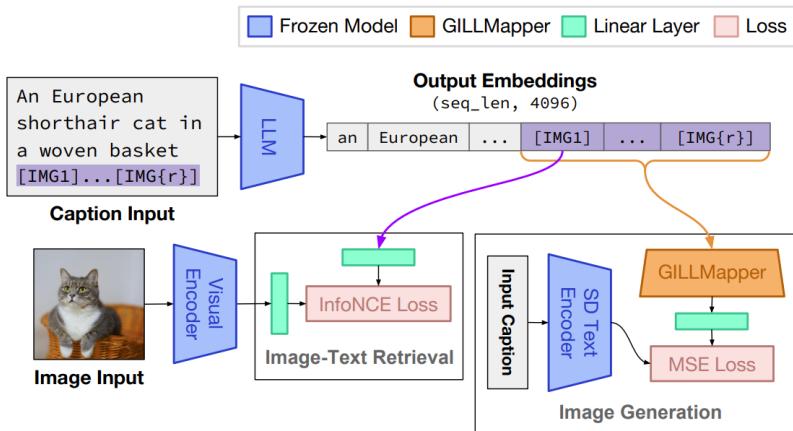
- [1] Generating Images with Multimodal Language Models. 2023
- [2] Generative Pretraining in Multimodality. 2023
- [3] MiniGPT-5: Interleaved Vision-and-Language Generation via Generative Vokens. 2023
- [4] DreamLLM: Synergistic Multimodal Comprehension and Creation. 2023
- [5] LLaVA-Plus: Learning to Use Tools for Creating Multimodal Agents. 2023

...

Unified MLLM: Perceiving + Generation

- Image

- GILL

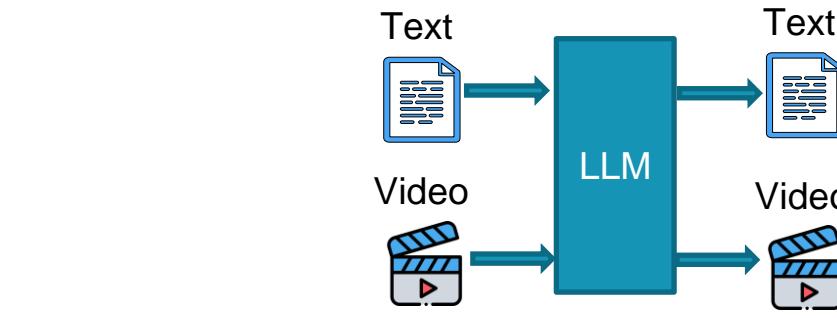


- [1] Generating Images with Multimodal Language Models. 2023
[2] Generative Pretraining in Multimodality. 2023

* Unified MLLM: Perceiving + Generation

- Video

- + GPT4Video
- + VideoPoet
- + Video-LaVIT
- + ...



Central LLMs take as input both texts and videos, after semantics comprehension, and generate both texts and videos.

[1] GPT4Video: A Unified Multimodal Large Language Model for Instruction-Followed Understanding and Safety-Aware Generation. 2023

[2] VideoPoet: A Large Language Model for Zero-Shot Video Generation. 2023

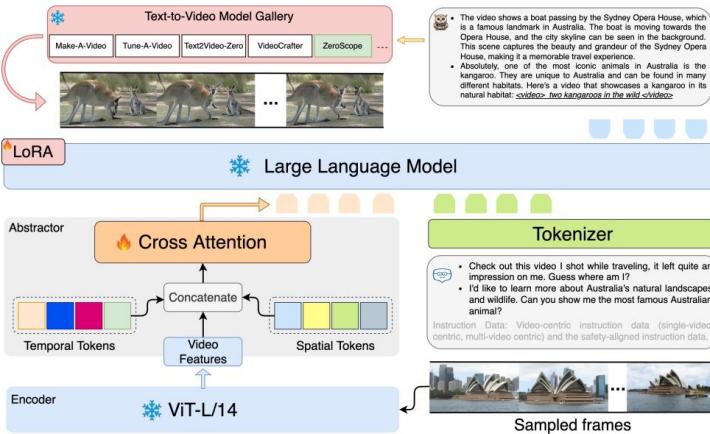
[3] Video-LaVIT: Unified Video-Language Pre-training with Decoupled Visual-Motional Tokenization. 2024

...

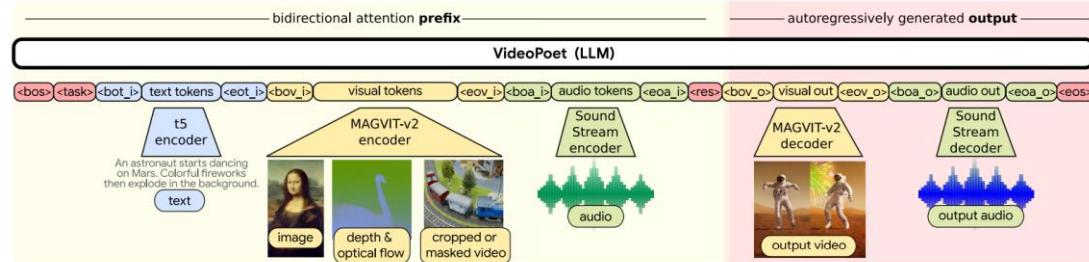
Unified MLLM: Perceiving + Generation

• Video

+ GPT4Video



+ VideoPoet



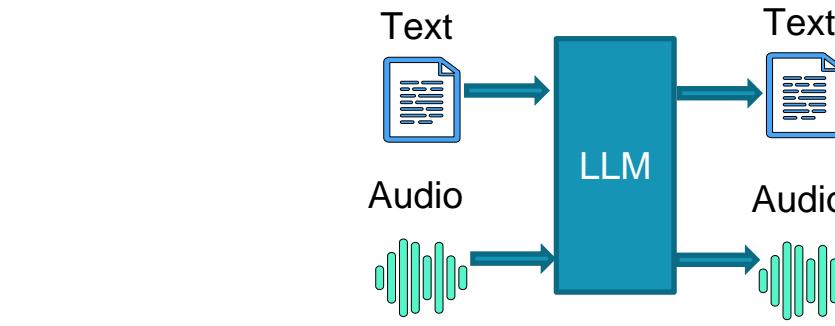
[1] GPT4Video: A Unified Multimodal Large Language Model for Instruction-Followed Understanding and Safety-Aware Generation. 2023

[2] VideoPoet: A Large Language Model for Zero-Shot Video Generation. 2023

* Unified MLLM: Perceiving + Generation

- **Audio**

- + AudioGPT,
- + SpeechGPT,
- + VIOLA,
- + AudioPaLM,
- + ...



Central LLMs take as input both texts and audio, after semantics comprehension, and generate both texts and audio.

[1] *AudioGPT: Understanding and Generating Speech, Music, Sound, and Talking Head.* 2023

[2] *SpeechGPT: Empowering Large Language Models with Intrinsic Cross-Modal Conversational Abilities.* 2023

[3] *VioLA: Unified Codec Language Models for Speech Recognition, Synthesis, and Translation.* 2023

[4] *AudioPaLM: A Large Language Model That Can Speak and Listen.* 2023

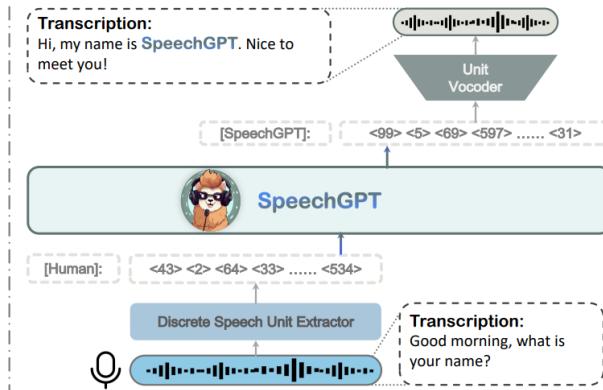
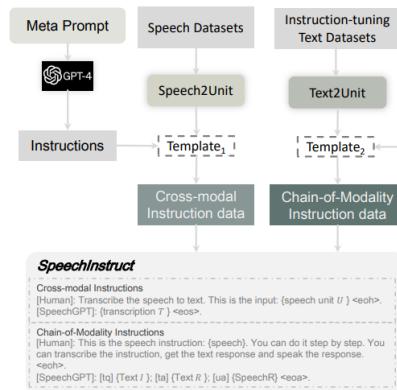
...

Unified MLLM: Perceiving + Generation

• Audio

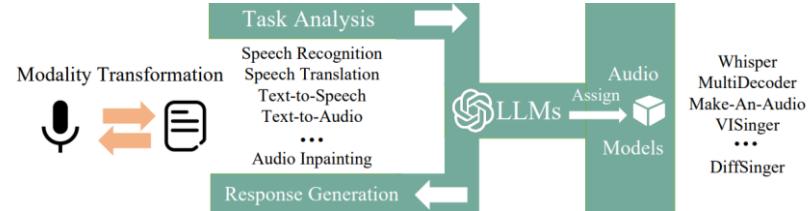
+

SpeechGPT



+

AudioGPT



[1] *SpeechGPT: Empowering Large Language Models with Intrinsic Cross-Modal Conversational Abilities*. 2023

[2] *AudioGPT: Understanding and Generating Speech, Music, Sound, and Talking Head*. 2023

* Unified MLLM: Harnessing Multi-Modalities

- Scenarios:



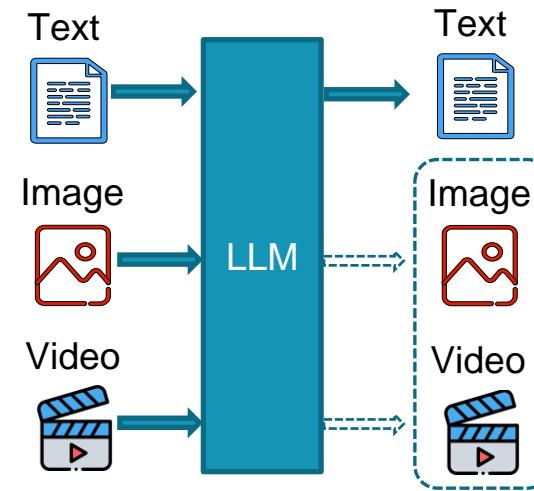
*In reality, modalities often have strong interconnections simultaneously. Thus, it is frequently necessary for MLLMs to handle the understanding of **multiple non-textual modalities at once**, rather than just one single (non-textual) modality.*

- + Image+Video
- + Audio+Video
- + Image+Video+Audio
- + Any-to-Any
- + ...

* Unified MLLM: Harnessing Multi-Modalities

- **Text+Image+Video**

- + Video-LLaVA
- + Chat-UniVi
- + LLaMA-VID
- + ...



Central LLMs take as input texts, image and video, after semantics comprehension, and generate texts (maybe also image and video, or combination).

[1] Video-LLaVA: Learning United Visual Representation by Alignment Before Projection. 2023

[2] Chat-UniVi: Unified Visual Representation Empowers Large Language Models with Image and Video Understanding. 2023

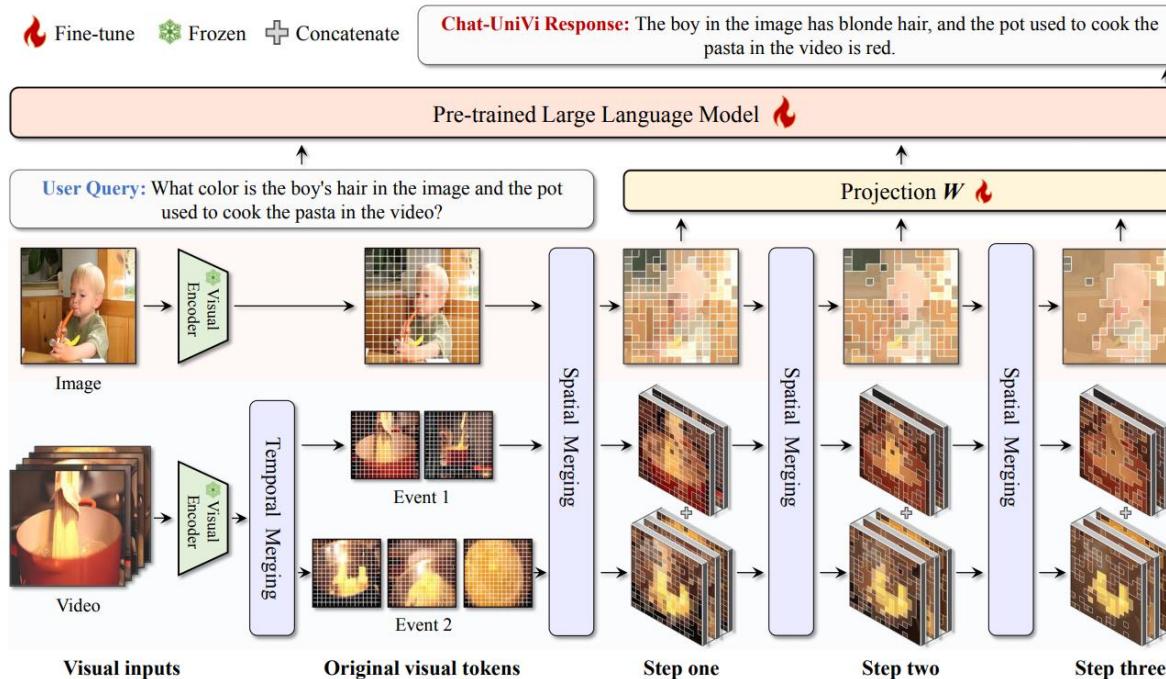
[3] LLaMA-VID: An Image is Worth 2 Tokens in Large Language Models. 2023

...

* Unified MLLM: Harnessing Multi-Modalities

- Text+Image+Video

+ Chat-UniVi

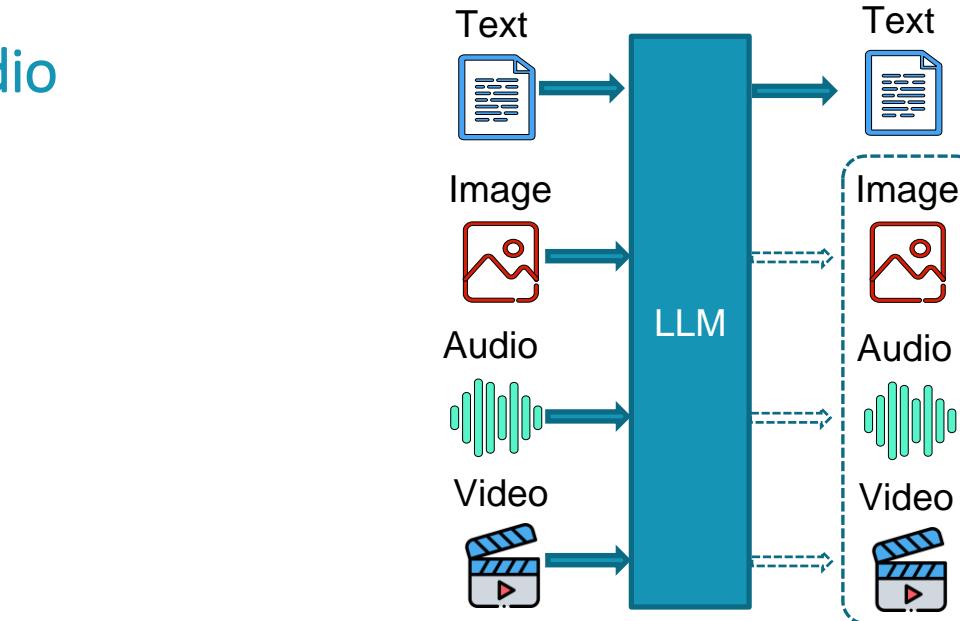


[1] Chat-UniVi: Unified Visual Representation Empowers Large Language Models with Image and Video Understanding. 2023

* Unified MLLM: Harnessing Multi-Modalities

- **Text+Image+Video+Audio**

- + Panda-GPT
- + Video-LLaMA
- + AnyMAL
- + Macaw-LLM
- + VideoPoet
- + ImageBind-LLM
- + LLMBind
- + LLaMA-Adapter
- + ...



 *Central LLMs take as input texts, audio, image and video, and generate texts (maybe also audio, image and video, or combination).*

[1] PandaGPT: One Model to Instruction-Follow Them All. 2023

[2] Video-LLaMA: An Instruction-tuned Audio-Visual Language Model for Video Understanding. 2023

[3] AnyMAL: An Efficient and Scalable Any-Modality Augmented Language Model. 2023

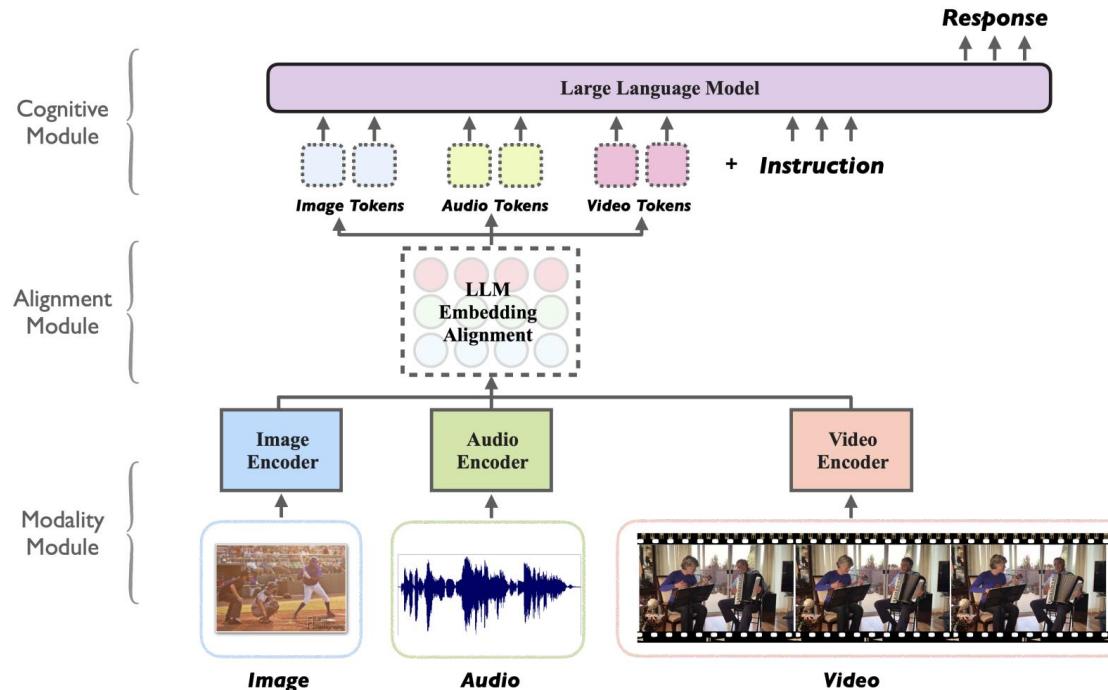
[4] Macaw-LLM: Multi-Modal Language Modeling with Image, Audio, Video, and Text Integration. 2023

...

* Unified MLLM: Harnessing Multi-Modalities

- Text+Image+Video+Audio

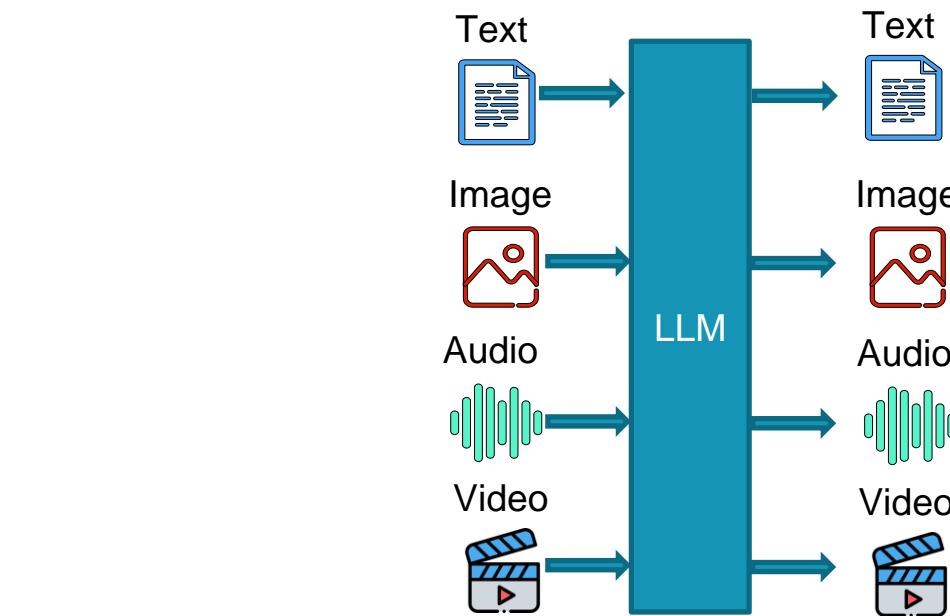
+ Macaw-LLM



* Unified MLLM: Harnessing Multi-Modalities

- Any-to-Any MLLM

- + NExT-GPT
- + Unified-IO 2 (w/o video)
- + AnyGPT (w/o video)
- + CoDi-2
- + Modaverse
- + ...



Central LLMs take as input texts, audio, image and video, and freely generate texts, audio, image and video, or combination.

[1] NExT-GPT: Any-to-Any Multimodal LLM. 2023

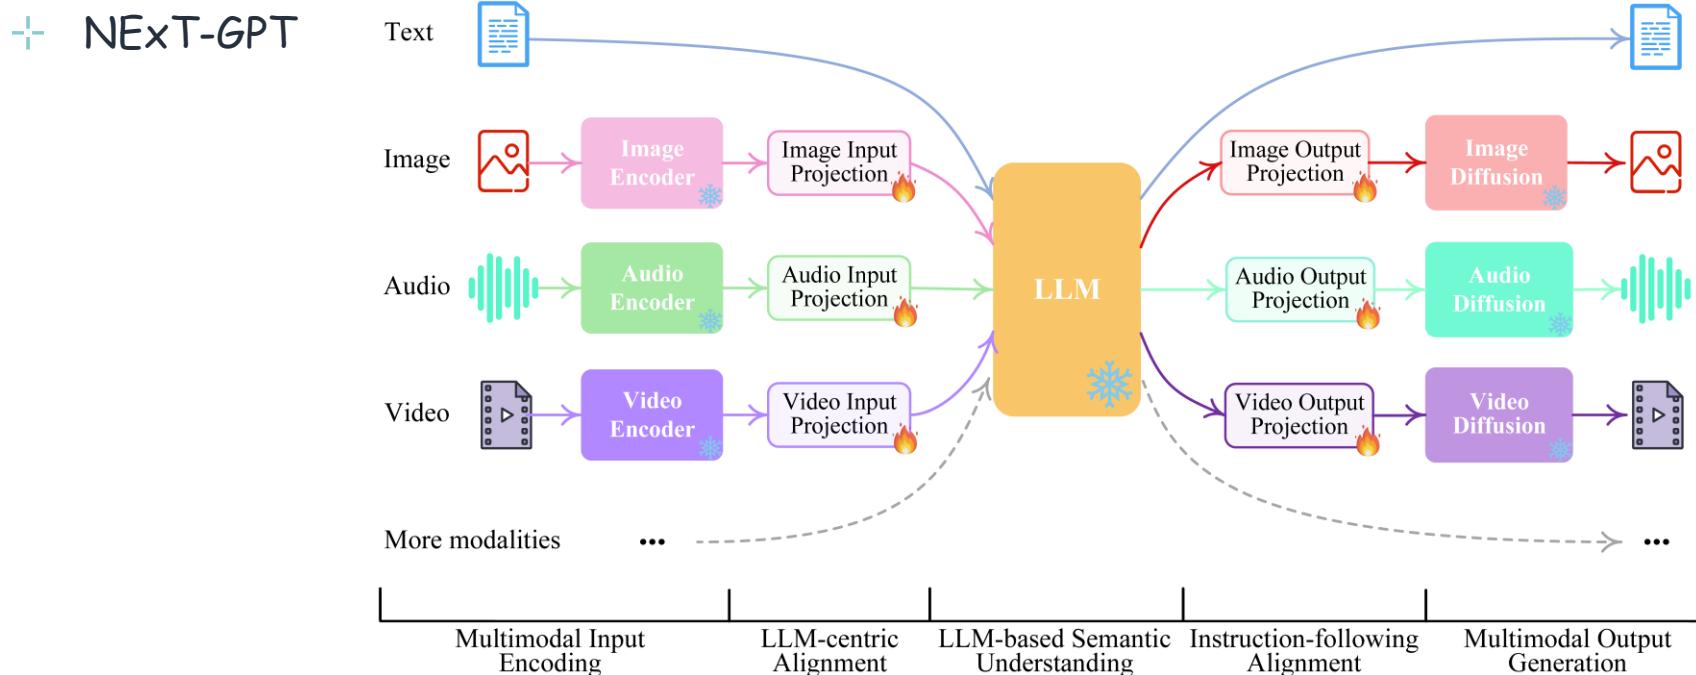
[2] AnyGPT: Unified Multimodal LLM with Discrete Sequence Modeling. 2023

[3] CoDi-2: In-Context, Interleaved, and Interactive Any-to-Any Generation. 2023

[4] Modaverse: Efficiently Transforming Modalities with LLMs. 2023

* Unified MLLM: Harnessing Multi-Modalities

- Any-to-Any MLLM



* Unified MLLM: Harnessing Multi-Modalities

- Any-to-Any MLLM NExT-GPT
- + NExT-GPT



Text + Audio
↓
Text + Image + Video

Project: <https://next-gpt.github.io>

Paper: <https://arxiv.org/pdf/2309.05519.pdf>

Code: <https://github.com/NExT-GPT/NExT-GPT>

* Fine-grained Capability of MLLM

• Pixel-level Vision MLLM



The vision MLLMs described above generally only support coarse-grained, instance-level visual understanding. This can lead to **imprecise visual interpretations**. Also due to the lack of visual grounding, these MLLMs will potentially **produce hallucinations**.

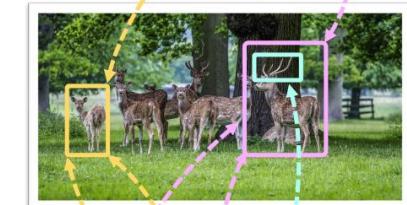
- + Visual Grounding
- + Visual Segmentation
- + Visual Editing
- + Visual Inpainting
- + ...



Describe the image in detail. Please output interleaved segmentation mask



What is the difference between this deer and another deer?



The main difference between the two deer is their size and antlers. The adult deer is larger with more prominent antlers, while the younger deer is smaller with no visible antlers.

Shikra



The image showcases a large, white building with a red roof, surrounded by a well-manicured lawn and palm trees. The sky is visible over the building, the pavement, and the grass. The grass is also seen extending to the pavement.

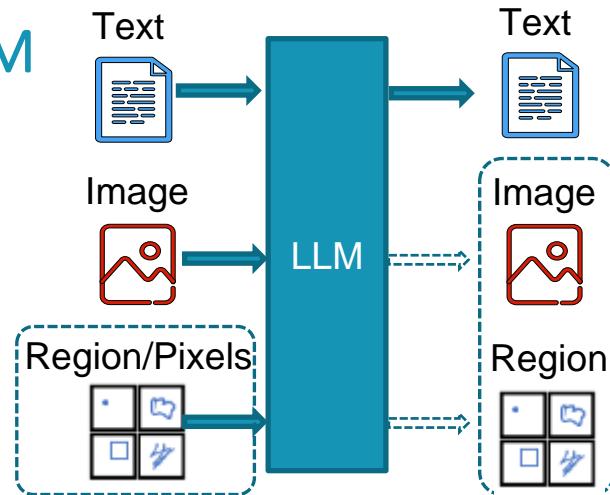
Fine-grained Capability of MLLM

- **Image-oriented Pixel-wise Regional MLLM**

- + GPT4RoI
- + NExT-Chat
- + MiniGPT-v2
- + Shikra
- + Kosmos-2
- + GLaMM
- + LISA
- + DetGPT
- + Osprey
- + PixelLM
- + LION
- + ...



Users input an image (potentially specifying a region), and the LLM outputs content based on its understanding, grounding the visual content to specific pixel-level regions of the image.



- [1] GPT4RoI: Instruction Tuning Large Language Model on Region-of-Interest. 2023
- [2] NExT-Chat: An LMM for Chat, Detection and Segmentation. 2023
- [3] MiniGPT-v2: large language model as a unified interface for vision-language multi-task learning. 2023
- [4] Osprey: Pixel Understanding with Visual Instruction Tuning. 2023
- [5] GLaMM: Pixel Grounding Large Multimodal Model. 2023
- [6] Kosmos-2: Grounding Multimodal Large Language Models to the World. 2023
- [7] DetGPT: Detect What You Need via Reasoning. 2023
- [8] PixelLM: Pixel Reasoning with Large Multimodal Model. 2023
- [9] Lisa: Reasoning segmentation via large language model. 2023
- [10] Shikra: Unleashing Multimodal LLM's Referential Dialogue Magic. 2023

...

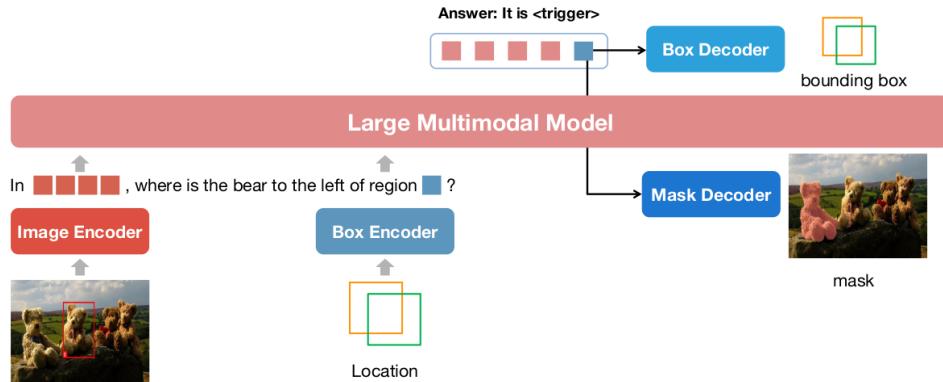
✳️ Overview of Modality and Functionality

	Modality (w/ Language)			
	Image	Video	Audio	3D
Input-side Perceiving	Flamingo, Kosmos-1, Blip2, mPLUG-Owl, Mini-GPT4, LLaVA, InstructBLIP, VPGTrans, CogVLM, Monkey, Chameleon, Otter, Qwen-VL, GPT-4v, SPHINX, Yi-VL, Fuyu, ...	VideoChat, Video-ChatGPT, Video-LLaMA, PandaGPT, MovieChat, Video-LLaVA, LLaMA-VID, Momentor, ...	AudioGPT, SpeechGPT, VIOLA, AudioPaLM, SALMONN, MU-LLaMA, ...	3D-LLM, 3D-GPT, LL3DA, SpatialVLM, PointLLM, PointBind, ...
	[Pixel-wise] GPT4RoI, LION, MiniGPT-v2, NExT-Chat, Kosmos-2, GLaMM, LISA, DetGPT, Osprey, PixelLM, ...	[Pixel-wise] PG-Video-LLaVA, Merlin, MotionEpic, ...	-	-
	Video-LLaVA, Chat-UniVi, LLaMA-VID		-	-
	Panda-GPT, Video-LLaMA, AnyMAL, Macaw-LLM, Gemini, VideoPoet, ImageBind-LLM, LLMBind, LLaMA-Adapter, ...			-
Perceiving + Generating	GILL, EMU, MiniGPT-5, DreamLLM, LLaVA-Plus, InternLM-XComposer2, SEED-LLaMA, LaVIT, Mini-Gemini, ...	GPT4Video, Video-LaVIT, VideoPoet, ...	AudioGPT, SpeechGPT, VIOLA, AudioPaLM, ...	-
	[Pixel-wise] Vitron		-	-
	NExT-GPT, Unified-IO 2, AnyGPT, CoDi-2, Modaverse, ViT-Lens, ...			-

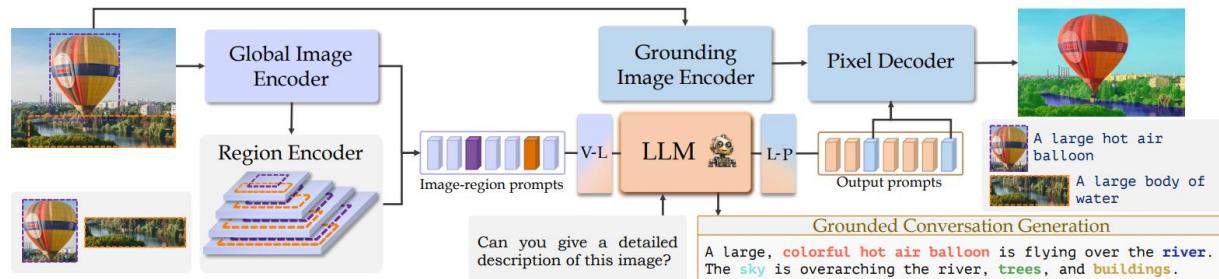
Fine-grained Capability of MLLM

- Image-oriented Pixel-wise

+ NExT-Chat



+ GLaMM



Referring Expression Segmentation
Can you segment the balloon in this given picture?
Sure, the segmentation is [image of segmented balloon]

Image Level Captioning
Please describe the given image. A hot air balloon is seen in the air, flying over a river. The balloon, high up in the sky, is yellow and red in color, with a bridge and trees visible in the background.

Region Level Captioning
Could you tell me more about this region? A yellow, red and blue hot air balloon

Phrase Grounding
Can you segment phrases in, "A <hot air balloon> over the <river>" A hot air balloon over the <river>

* Fine-grained Capability of MLLM

- **Image-oriented Pixel-wise Regional MLLM**



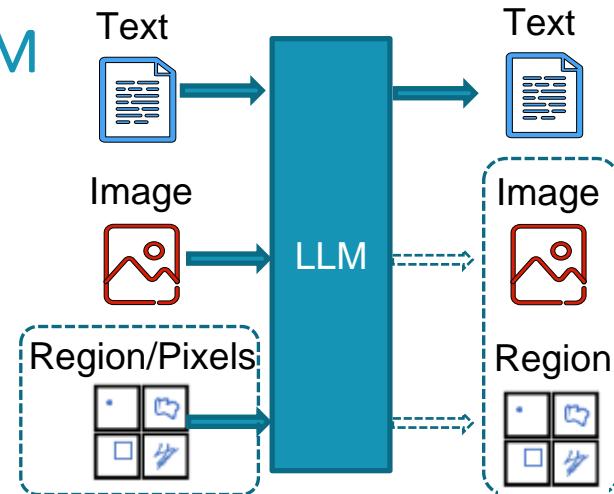
Pixel-level Awareness at Input/Output

+ Output-side Only Pixel-wise Awareness

LISA, PixelLM, DetGPT, MiniGPT-v2, LION

+ Input-&Output-side Pixel-wise Awareness

NExT-Chat, GPT4RoI, Shikra,
KOSMOS-2, GLaMM, Osprey



* Fine-grained Capability of MLLM

- **Image-oriented Pixel-wise Regional MLLM**

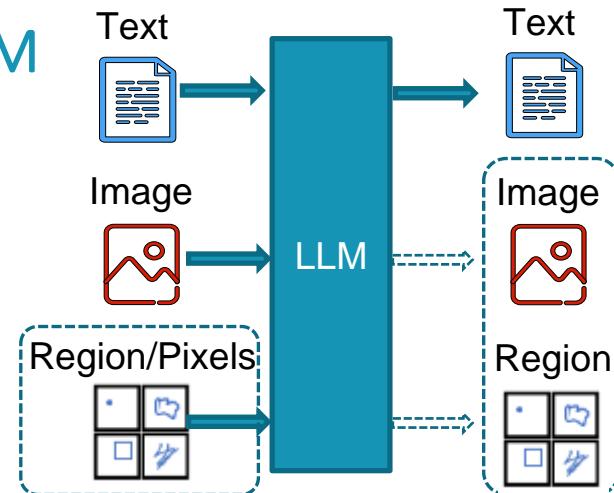
-  **Pixel Granularity**

-  **Bounding-box Coordinates**

- NExT-Chat, GPT4RoI, Shikra, LION, KOSMOS-2, DetGPT, MiniGPT-v2

-  **Finer-grained Mask-based Segments**

- NExT-Chat, LISA, PixelLM, GLaMM, Osprey



* Fine-grained Capability of MLLM

- **Image-oriented Pixel-wise Regional MLLM**



User Input Interaction

- + No Image User Interaction

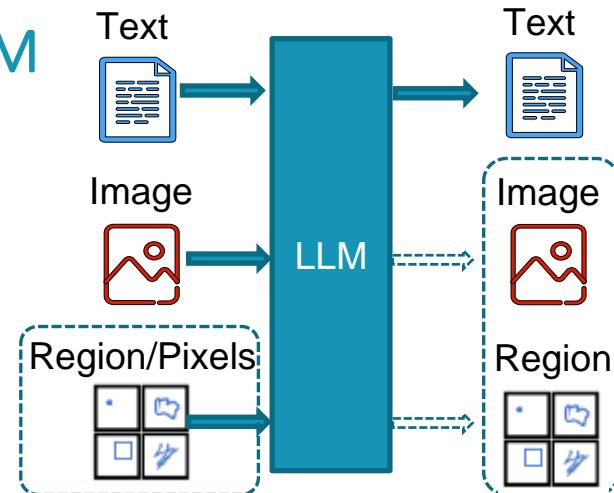
LISA, PixelLM, DetGPT, MiniGPT-v2, LION

- + Bounding-box Coordinates

GPT4RoI, Shikra, KOSMOS-2, GLaMM

- + User Sketches

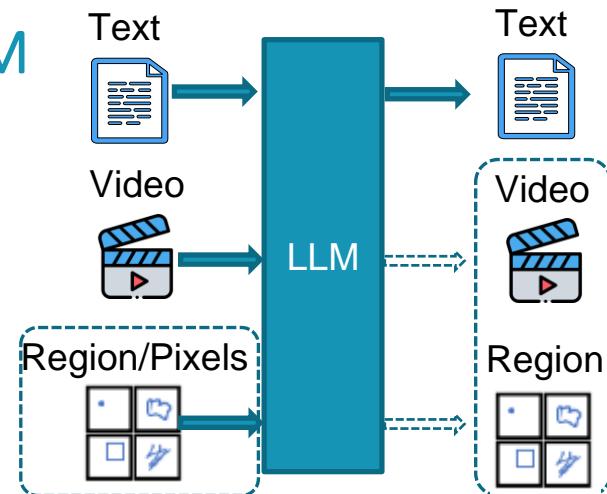
NExT-Chat, Osprey,



* Fine-grained Capability of MLLM

- Video-oriented Pixel-wise Regional MLLM

- + PG-Video-LLaVA
- + Merlin
- + MotionEpic
- + ...



Users input an video (potentially specifying a region), and the LLM outputs content based on its understanding, grounding or tracking the content to specific pixel-level regions of the video.

[1] PG-Video-LLaVA: Pixel Grounding in Large Multimodal Video Models. 2023

[2] Merlin: Empowering Multimodal LLMs with Foresight Minds. 2023

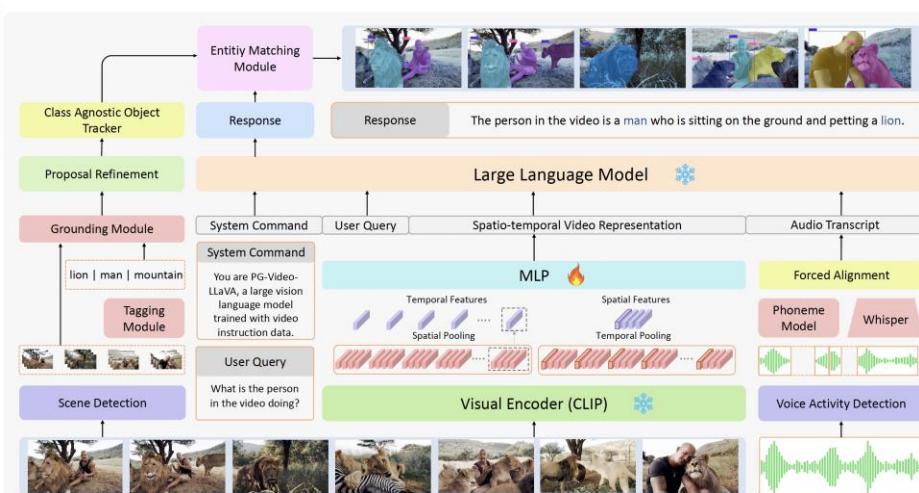
[3] Video-of-Thought: Step-by-Step Video Reasoning from Perception to Cognition. 2024

...

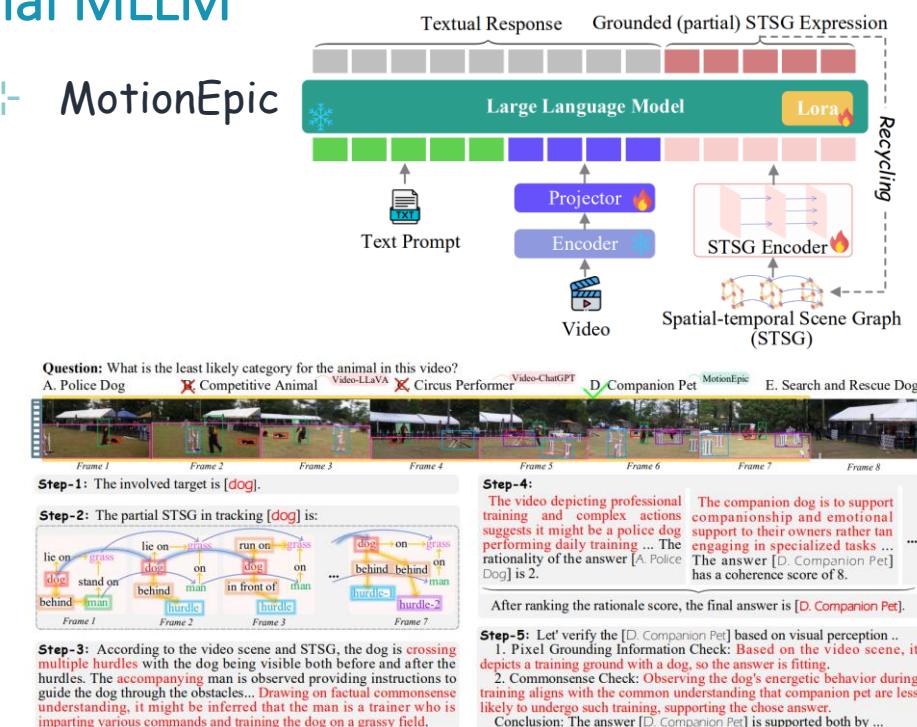
Fine-grained Capability of MLLM

- Video-oriented Pixel-wise Regional MLLM

- + PG-Video-LLaVA



- + MotionEpic



- [1] PG-Video-LLaVA: Pixel Grounding in Large Multimodal Video Models. 2023
[2] Video-of-Thought: Step-by-Step Video Reasoning from Perception to Cognition. 2024

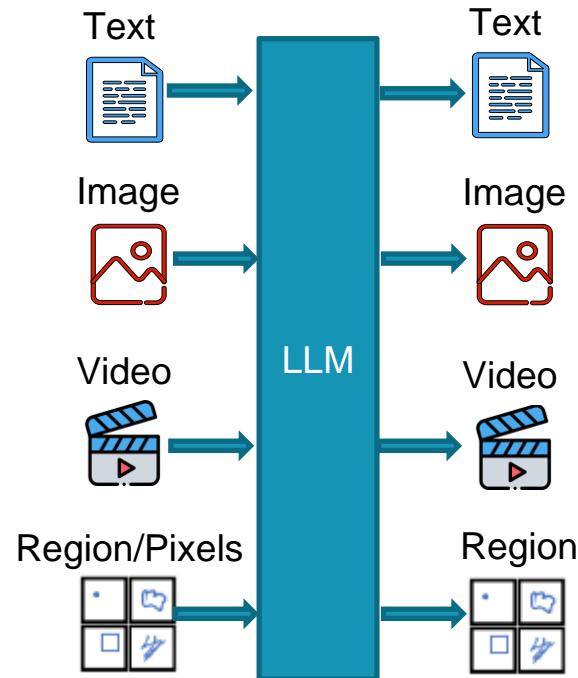
* Fine-grained Capability of MLLM

- Unified Pixel-wise MLLM

- + Vitron



Users input either an image or video (potentially specifying a region), and the LLM outputs content based on its understanding, generating, grounding or tracking the content to specific pixel-level regions of the image, video.



Fine-grained Capability of MLLM

Unified

Vitron

Model	Vision Supporting		Pixel/Regional Understanding	Segmenting/Grounding	Generating	Editing
	Image	Video				
Flamingo [1]	✓	✗	✗	✗	✗	✗
BLIP-2 [45]	✓	✗	✗	✗	✗	✗
MiniGPT-4 [126]	✓	✗	✗	✗	✗	✗
LLaVA [57]	✓	✗	✗	✗	✗	✗
GILL [39]	✓	✗	✗	✗	✓	✗
Emu [90]	✓	✗	✗	✗	✓	✗
MiniGPT-5 [125]	✓	✗	✗	✗	✓	✗
DreamLLM [23]	✓	✗	✗	✗	✓	✗
GPT4RoI [122]	✓	✗	✓	✓	✗	✗
NExT-Chat [118]	✓	✗	✓	✓	✗	✗
MiniGPT-v2 [13]	✓	✗	✓	✓	✗	✗
Shikra [14]	✓	✗	✓	✓	✗	✗
Kosmos-2 [72]	✓	✗	✓	✓	✗	✗
GLaMM [78]	✓	✗	✓	✓	✗	✗
Osprey [117]	✓	✗	✓	✓	✗	✗
PixelLM [79]	✓	✗	✓	✓	✗	✗
LLaVA-Plus [58]	✓	✗	✗	✓	✓	✓
VideoChat [46]	✗	✓	✗	✗	✗	✗
Video-LLaMA [120]	✗	✓	✗	✗	✗	✗
Video-LLaVA [52]	✓	✓	✗	✗	✗	✗
Video-ChatGPT [61]	✗	✓	✗	✗	✗	✗
GPT4Video [99]	✗	✓	✗	✗	✓	✗
PG-Video-LLaVA [67]	✗	✓	✓	✓	✗	✗
NExT-GPT [104]	✓	✓	✗	✗	✓	✗
VITRON (Ours)	✓	✓	✓	✓	✓	✓

[1] VITRON: A Unified Pixel-level Vision LLM for Understanding, Generating, Segmenting, Editing. 2024

Fine-

- Unified
- + Vitron

Visual Understanding ↑ Visual Generating ↓

Low-level Visual Semantics

Vision Segmentation & Grounding



Phrase Grounding



Video Object Segmentation (Tracking)



Video Grounding



High-level Visual Semantics

Pixel-level Vision Understanding

Image/Video Captioning



a drawing of a pink and blue pokemon

Referring Captioning



woman on the right in white shirt

Image QA



Q: Where is the child sitting?
A: arms

Video QA



Q: Why is the child sitting on a cardboard? A: trying to slide down the slope.

Language-Image Retrieval



panda 0.92
lion 0.85
ostrich 0.79
emu 0.70

Language-Video Retrieval



Video Temporal Grounding



Vision Synthesis & Generation

Text-to-Image Generation

A corgi runs on the grass



Text-to-Video Generation

A polar bear swimming.



Image-to-Video Generation



VITRON

Vision Editing & Inpainting

Adding



Removing



Replacing



Moving



Style Changing



Color Changing



✳️ Overview of Modality and Functionality

	Modality (w/ Language)			
	Image	Video	Audio	3D
Input-side Perceiving	Flamingo, Kosmos-1, Blip2, mPLUG-Owl, Mini-GPT4, LLaVA, InstructBLIP, VPGTrans, CogVLM, Monkey, Chameleon, Otter, Qwen-VL, GPT-4v, SPHINX, Yi-VL, Fuyu, ...	VideoChat, Video-ChatGPT, Video-LLaMA, PandaGPT, MovieChat, Video-LLaVA, LLaMA-VID, Momentor, ...	AudioGPT, SpeechGPT, VIOLA, AudioPaLM, SALMONN, MU-LLaMA, ...	3D-LLM, 3D-GPT, LL3DA, SpatialVLM, PointLLM, PointBind, ...
	[Pixel-wise] GPT4RoI, LION, MiniGPT-v2, NExT-Chat, Kosmos-2, GLaMM, LISA, DetGPT, Osprey, PixelLM, ...	[Pixel-wise] PG-Video-LLaVA, Merlin, MotionEpic, ...	-	-
	Video-LLaVA, Chat-UniVi, LLaMA-VID		-	-
	Panda-GPT, Video-LLaMA, AnyMAL, Macaw-LLM, Gemini, VideoPoet, ImageBind-LLM, LLMBind, LLaMA-Adapter, ...			-
Perceiving + Generating	GILL, EMU, MiniGPT-5, DreamLLM, LLaVA-Plus, InternLM-XComposer2, SEED-LLaMA, LaVIT, Mini-Gemini, ...	GPT4Video, Video-LaVIT, VideoPoet, ...	AudioGPT, SpeechGPT, VIOLA, AudioPaLM, ...	-
	[Pixel-wise] Vitron		-	-
	NExT-GPT, Unified-IO 2, AnyGPT, CoDi-2, Modaverse, ViT-Lens, ...			-

Fine-grained Capability of MLLM

- Unified Pixel-wise MLLM

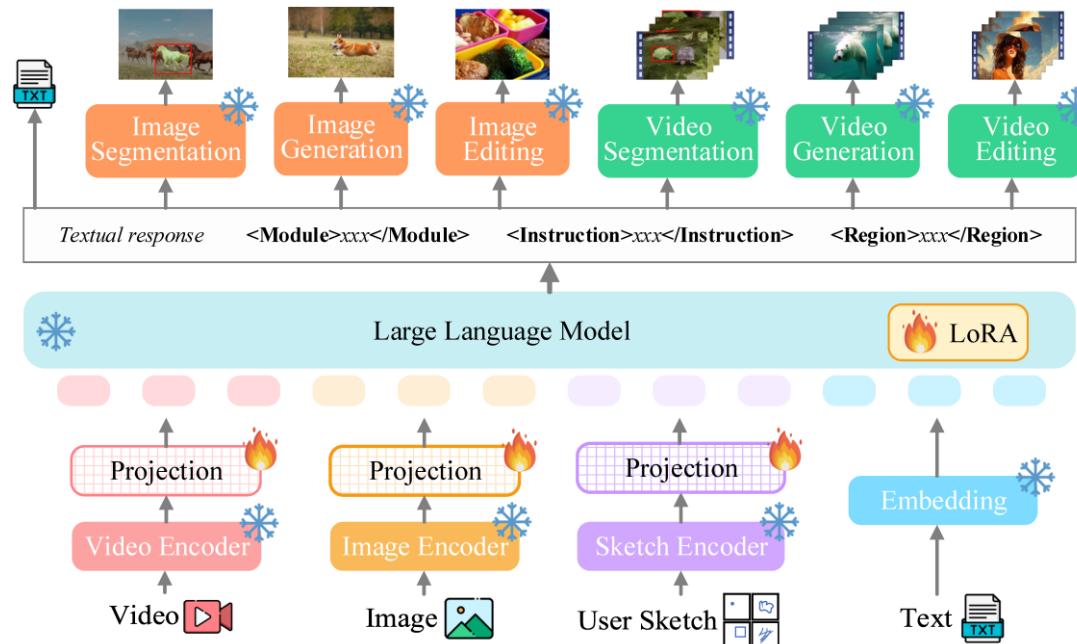
- + Vitron

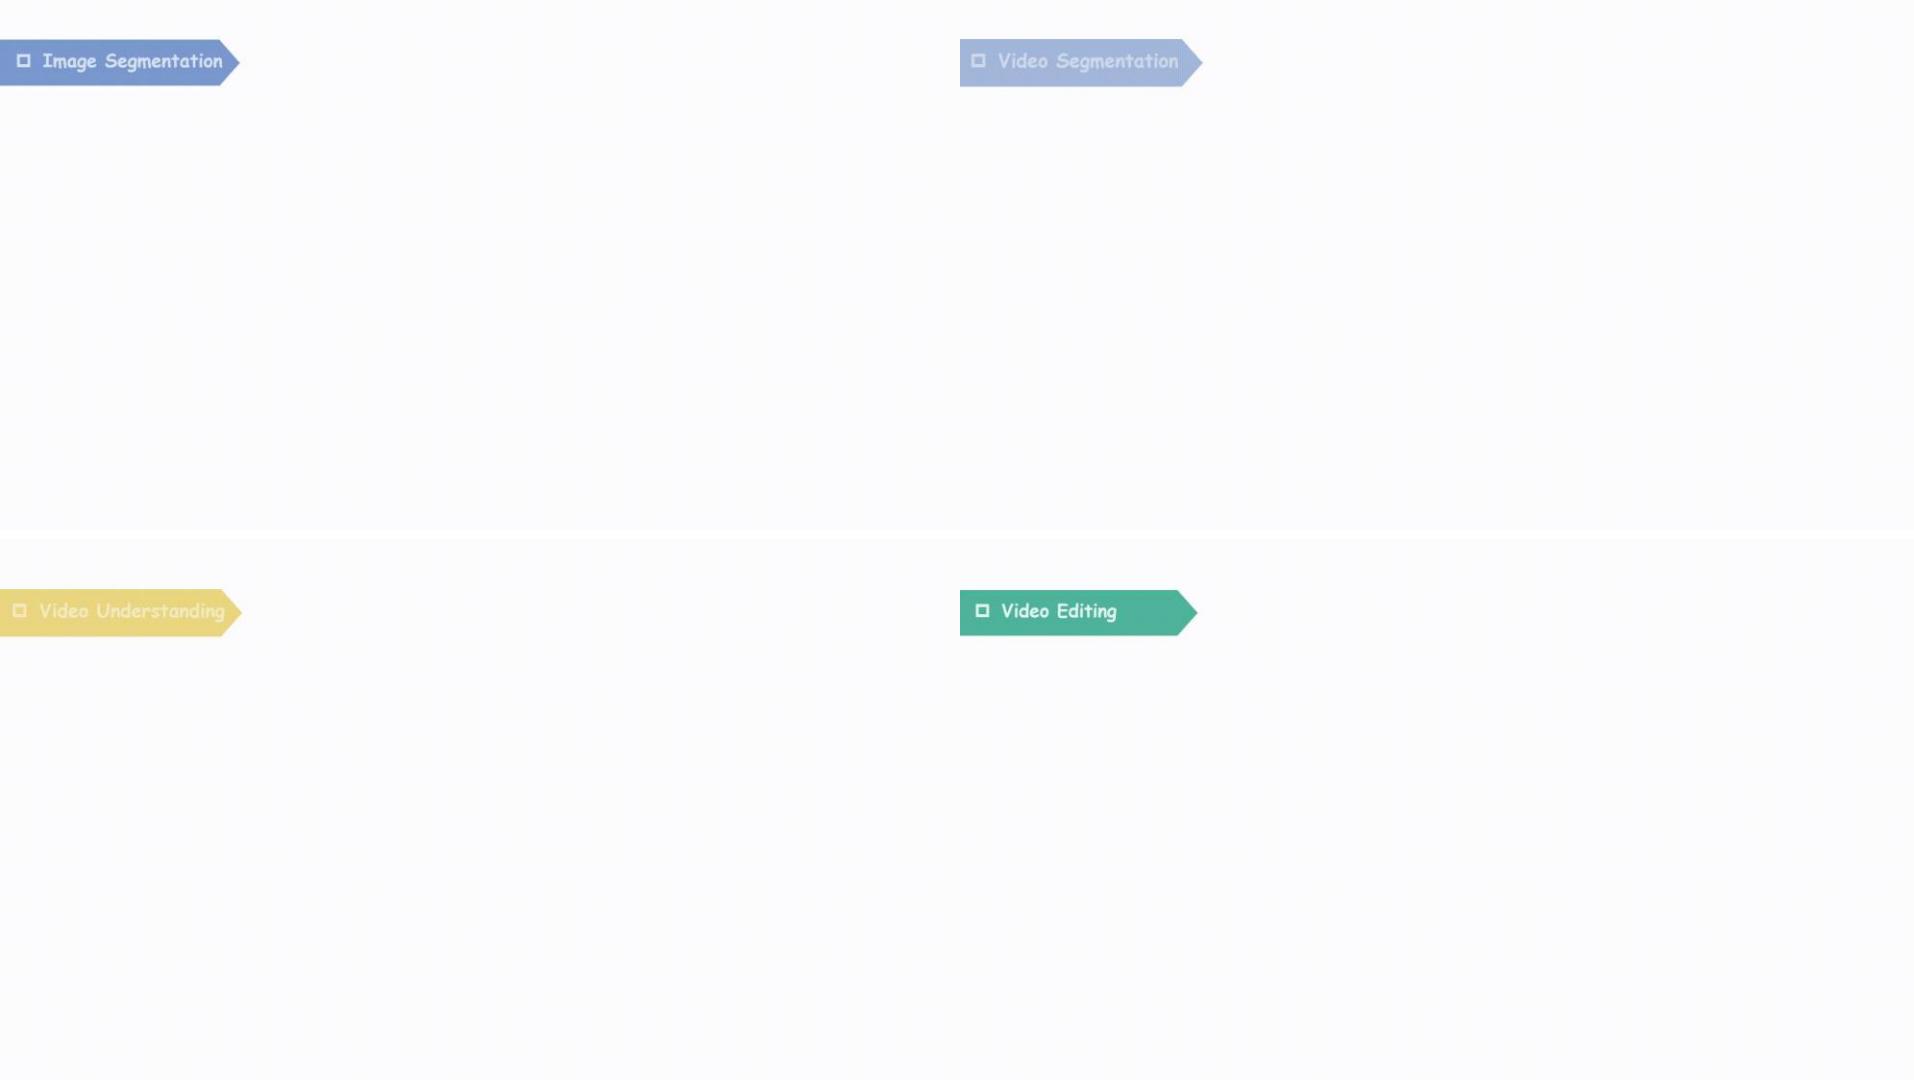


Project: <https://vitron-llm.github.io/>

Paper: <https://is.gd/aGu0VV>

Code&Demo: <https://github.com/SkyworkAI/Vitron>





✳ What's Next

- Angle-I: Unification of as Many Modalities & Tasks as Possible
 - + Modality Perspective: Going Broader

 *Currently, the majority of MLLM research focuses primarily on the integration of visual signals (e.g., Image, Video).*

✳ What's Next from Multimodal LLM to AGI

- Angle-I: Unification of as Many Modalities & Tasks as Possible

- + Modality Perspective: Going Broader

- Modalities in current NExT-GPT:

language 

image 

sound 

video 

- More modalities to go:

heat map 

code 

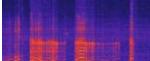
time series 

touch 

depth&3D 

infrared/radar 

document/table 

spectrogram 

smell 

graph 

✳ What's Next from Multimodal LLM to AGI

- Angle-I: Unification of as Many Modalities & Tasks as Possible

- + Task Perspective: Going Deeper

👉 Vision-based MLLM, *Vitron*, has focused on unifying image and video processing under the scope of pixel-wise tasks, ranging from low-level to high-level.

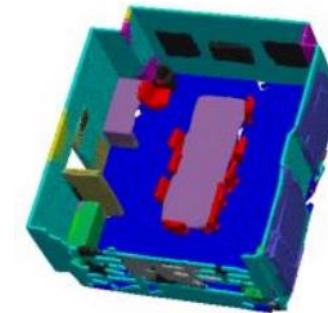
👉 The next step could involve expanding MLLM support on the task level to more in-depth levels.



Referring Segmentation



Panoptic Segmentation



3D Scene Segmentation

✳ What's Next from Multimodal LLM to AGI

- Angle-II: Stronger Generation Ability via Better Tokenization

- + Core Idea

 *High-quality multimodal generation requires the system to **recover a sufficient amount of detailed multimodal information from the core LLM.***

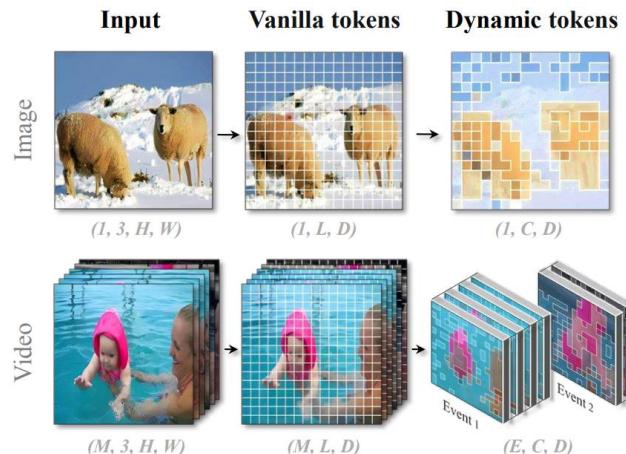
- + Remove the equivalence constraint between pre-LLM and post-LLM, as the roles of input and output multimodal tokens differ.
 - + Increase the information content of multimodal tokens to include more high-frequency details.

✳ What's Next from Multimodal LLM to AGI

- Angle-II: Stronger Generation Ability via Better Tokenization

- ✚ A Hot Trend: Video tokenization

👉 *Supporting both images and videos: more carefully model the spatial aspects of images and the temporal dynamics of videos.*



[1] LLaMA-VID: An Image is Worth 2 Tokens in Large Language Models. 2024

[2] Chat-UniVi: Unified Visual Representation Empowers Large Language Models with Image and Video Understanding. 2024

[3] Video-LaVIT: Unified Video-Language Pre-training with Decoupled Visual-Motional Tokenization. 2024

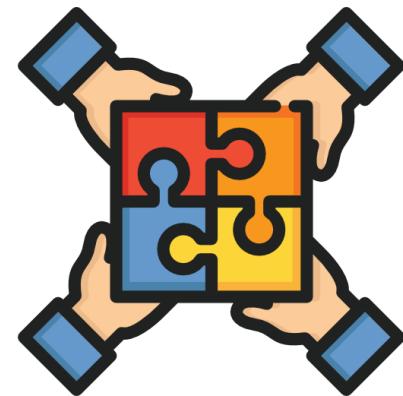
✳ What's Next from Multimodal LLM to AGI

- Angle-III: More Multimodality & Multi-Task Synergy

- + Core Idea

👉 Achieving a stronger MLLM, and potentially reaching AGI, necessitates enhanced Multimodality & Multi-Task Synergy for the MLLM generalist.

👉 Master *abductive reasoning* to facilitate *analogical thinking*, allowing different modalities and tasks, as well as the comprehension and generation processes, to mutually assist each other and create synergistic effects.



[1] Abductive reasoning: Logic, visual thinking, and coherence. 1997.

[2] Reasoning. <https://www.butte.edu/departments/cas/tipsheets/thinking/reasoning.html>

Thanks!

Any questions?

